

Coding Interview: Compressed Tries

Tries are data structures which can be used to implement symbol tables, much like hash tables can. If we define the alphabet for valid keys to be Σ , then a key is just a string on $s \in \Sigma^*$. As all symbol tables, tries must support the usual operations; e.g., in Python:

```
1 class Trie(Generic[V]):  
2     def search[V](self, key: str) -> Optional[V]:  
3         ...  
4     def insert[V](self, key: str, value: V):  
5         ...  
6     def delete[V](self, key: str):  
7         ...
```

Regular tries are m -ary trees which store prefixes in internal nodes and content on both leaves and internal nodes^[1]. Each node can have as many children as there are letters in Σ ; i.e., $|\Sigma| = m$, and each edge corresponds to a character in Σ . For instance, calling the sequence:

```
1 | trie = Trie[int]()
2 |
3 | trie.insert('carbonated', 1)
4 | trie.insert('carbonic', 3)
5 | trie.insert('carbide', 5)
6 | trie.insert('car', 7)
7 | trie.insert('bible', 9)
8 | trie.insert('biblical', 11)
```

would produce a trie that looks like the one in Fig. 1. Nodes that contain values are highlighted in blue, with the value on the side.




Figure 1. A regular trie.

Performing a search amounts to following the path corresponding to each letter in the key and testing: *i)* if we can get to the end of the key before bumping into a leaf, meaning the key is not in the trie and; *ii)* the node where our search ends contains a value. If it does not, then the key is not in the trie.

When we perform a search for, say, "car", we follow a path from the root through the link labelled *c*, then through the link labelled *a*, then through the link labelled *r*. The node pointed to by this last link contains the value 7, so we return it.

Tries are space-inefficient: unique prefixes are represented by long paths formed by single nodes without values (e.g. "carb" and "bibl" in Fig. 1), so we would like to implement what is known as *compressed* trie, and sometimes a PATRICIA trie^[2]. In compressed tries, we compact all single-child nodes with no values into a prefix within the node itself. A compressed trie equivalent to the one in Fig. 1 would look like:




Figure 2. A compressed trie.

Note that edges are labeled by path as before, and nodes can still have as many children as there are letters in Σ , but one-way paths are compressed into string prefixes within the node itself. Your job will be to implement `insert` and `search` on a compressed trie, using the language of your choosing.

Example input:

```
1  trie = CompressedTrie()
2
3  trie.insert('carbonated', 2)
4  trie.insert('carbonic', 3)
5  trie.insert('carbide', 4)
6  trie.insert('car', 1)
7  trie.insert('bible', 5)
8  trie.insert('biblical', 6)
9
10 print(str(trie))
```

example output:

```
<<: None>>
  <<car: 1>>
    <<b: None>>
      <<ide: 4>>
        <<on: None>>
          <<ic: 3>>
            <<ated: 2>>
        <<bibl: None>>
          <<ical: 6>>
            <<e: 5>>
```

1. Depending on how you implement them - if you choose to a special string termination character like C strings, then you do not need to store values in internal nodes.

↵

2. Though PATRICIA tries are originally binary radix trees, the term is nowadays often used in reference to any type of compressed trie. ↵